

Data Mining in Healthcare

Orfeas Ilias Gkonis
Georgios Chrysovelidis
Nikolaos Laoutaris
Vasileios Kesopoulos

¹ MSc in Data Science

² **Univeristy**:International Hellenic University

³ **Location**:Thessaloniki,Greece

Abstract

Data generation in the healthcare sector is rising at a remarkable rate due to the widespread adoption of electronic health records, medical imaging, and wearable sensors. Many opportunities but also difficulties are presented by this increase in data volume and complexity for extracting valuable insights using data mining, which is a key area within data science and has become increasingly important in healthcare, as it offers the potential to discover essential patterns and trends in large datasets, assisting in early disease detection, clinical decision support, patient outcome prediction, hospital operation optimization and many more. This literature review provides a summary of current research in data mining applications for healthcare, focusing on methodologies, real-world applications, and ongoing challenges. Even though data mining offers great potential for revolutionizing healthcare, challenges such as data diversity, privacy issues, and the demand of model explainability continue to be important areas that need to be addressed to fully harness its capabilities.

Index Terms: Data Mining, Healthcare, Applications.

1 Introduction

With the adoption of electronic health records, medical imaging and wearable sensors, there is an unprecedented rise in data generation in the healthcare industry. The volume and complexity of these data presents both opportunities and challenges in gaining actionable insights.

Data mining is a key area of Data Science and refers to the process of discovering meaningful unseen patterns and trends in large datasets. Data mining techniques have become increasingly vital in the healthcare domain, as they detect diseases early, support clinical decision making, predict patient outcomes and optimize hospital operations.

Despite its potential, there are several challenges in the application of data mining in healthcare, including issues with data heterogeneity, privacy concerns, and the need for interpretable models that align with medical standards.

With this literature survey we aim to provide an overview of current research in the field of data mining applications to healthcare, highlighting methodologies, real world applications and ongoing challenges, with the goal to identify promising avenues for future exploration

2 Fundamental Concepts and Background

2.1 Data Mining

In the last decades, most of the organizations have been digitally transformed, and transitioned from paper based data to digital data. This was done as electronic systems provide significant benefits to both businesses and employees, by increasing the speed and efficiency of work and producing better results [1]. What is more, large amount of data is produced and saved on a daily basis, which raised caution at the beginning, since people were not sure about the feasibility of maintaining this increasing amount of data. However, it did not take long until they realized that invaluable information was hidden in this data. Businesses started

exploiting the data they have gathered throughout these years, to strengthen their operational efficiency, cut unnecessary costs and gain valuable insights. As a result, data mining became the center of attention, because of its ability to extract information from data. People have improved data mining algorithms and released newer versions in order to extract more robust and reliable information.

2.2 Healthcare Domain

All over the world, the healthcare domain is a booming sector in economy, and along with its growth come some challenges as well, like poor quality in services, higher costs, inefficiencies and increasing complexity [2]. For instance, expenditures in the U.S. healthcare between 2010 and 2015, increased from 2.6 trillion dollars to 3.2 trillion, an increase of 23%. From the 3.2 trillion dollars, around 21-47% of it consists of non-value added tasks like readmissions, inappropriate use of antibiotics, fraud, etc., or low quality care, as researchers found that around 251,454 patients in the U.S. die every year due to medical errors [3].

To address these challenges and improve the overall efficiency and effectiveness of healthcare delivery, advanced technologies and analytical methods are increasingly being adopted. One such promising approach is data mining, which enables healthcare providers to extract valuable insights from vast amounts of medical data.

The Healthcare domain can benefit from data mining techniques in order to provide better medical services to patients and in medical management decision making. Length of stay in a hospital, hospital ranking, effective treatments, fraud insurance claims by patients or providers, patient's readmission, drug recommendation systems, etc. are some of the services provided by data mining techniques[4]. The Centers of Medicare and Medicaid Services (CMS) used data mining to reduce hospital readmission rates and prevent 115 million dollars' worth of fraud from happening. What is more, there are estimations that each year in the U.S. healthcare system, 450 billion dollars can be saved by applying data mining

techniques [5].

Data regarding healthcare are collected from both hospitals – including patient’s data, diagnosis, drug prescriptions and treatment data – and healthcare practitioners, professionals who are licensed to provide medical services – including data collected from healthcare organizations like the World Health Organization (WHO) [6]. The data are collected from various sources, therefore they need to be organized, structured and processed before we get any meaningful results. Governmental or healthcare professional organizations provide the most reliable data.

3 Applications of Data Mining in Healthcare Support

3.1 Disease diagnosis and prognosis

Data mining can assist clinicians, due to its capabilities of identifying patterns, leading to early diagnosis of a vast number of different diseases, such as Alzheimer’s disease, dementia and different types of cancer as well. Through the deployment of machine learning methods upon mass electronic health records (HER), AI can assist doctors by improving accuracy and offering the potential to mitigate human errors often associated with experience-based diagnostic procedures [7].

Disease prognosis, that is, predicting the likely progression of a disease, is another application of data mining in the healthcare domain. It essentially involves a thorough analysis of historical patient data and their corresponding treatment outcomes to forecast future health trajectories for individuals. A notable example is the prediction of cancer patients’ survivability rates using neural networks, decision trees, and logistic regression [8].

3.2 Length of Stay (LOS) prediction

Hospital length of stay (LOS) prediction is an application of data mining in healthcare, with the aim of forecasting the duration of a patient’s hospitalization for the purpose of optimizing resource allocation (bed allocation) and for improving operational efficiency (avoiding over-stuffing). To make this possible, historical patient data are leveraged containing clinical measurements, admission types, comorbidities, diagnosis types and demographics as well [9]. Then, several data mining methods can be employed to identify patterns and estimate hospitalization duration. Moreover, it is important to note that research has indicated that emergency admissions and chronic conditions are frequently correlated with longer stays while elective procedures may result in shorter durations, and this highlights the importance of these features in predictive modeling [10]. Ultimately, by accurately predicting LOS, patient care is additionally improved by reducing wait times and preventing rushed discharges.

3.3 Drug-drug interactions (DDIs)

Advancements in natural language processing (NLP) have played a key role in the development of complex algorithms that can effectively identify drug-drug interactions in medical literature [11]. It is well established that taking multiple medications at the same time may induce negative effects of either medication, or raise, or even decrease its effects, and the danger increases with the number of drugs being taken. To overcome the challenge of monitoring every known drug-drug interaction, machine learning algorithms have been deployed to extract information from medical literature. These models can analyze molecular structures, reported adverse

events, metabolic pathways, target proteins and other features, and then flag any drug combinations that may cause unexpected side effects [12]. In essence, those predictive models allow clinicians to be alerted to potential DDIs before they cause any damage and thereby preventing negative consequences and leading to enhancement of patient safety.

3.4 Personalized medicine

Personalized medicine, also known as precision medicine, is a medical model that divides people into different categories for the purpose of tailoring medical decisions, practices and interventions to each patient based on their individual characteristics. So, instead of employing a universal strategy to treat the patient, it takes into account the unique person’s genetic composition, environment, lifestyle and other factors [13]. Personalized medicine can be approached by using data mining techniques with remarkable results. Through the examination of extensive and diverse datasets containing patient genomics, transcriptomics, proteomics, imaging data, and more, it is feasible to create advanced medical models and thereby optimizing healthcare outcomes, minimizing adverse effects and enhancing the efficiency of healthcare services [14]. Some known applications in this area are: drug response prediction, adverse drug event prediction, risk factors identification and treatment response prediction.

4 Methodologies and Algorithms

As we have seen, data mining has become a cornerstone of modern healthcare services. In this section we outline the key data mining techniques that make use of massive, often heterogeneous and unstructured data, linking them to specific types of healthcare support tasks and common goals. These methodologies range from classic statistical models to advanced machine learning and are applied to achieve a broad spectrum of results, such as risk stratification, behavioral analysis, care personalization and beyond.

4.1 Classification

These are among the most frequently used algorithms in data mining due to their massive utility in predictive modeling. Classification assigns categorical labels based on input features extracted from patient data. Decision Trees are very interpretable classification structures, which makes them popular in clinical decision making. Random Forests extend this model by using ensembles of individual trees that often achieve higher accuracy. Logistic Regression, a much simpler model, is still being used in epidemiology due to its being robust and interpretable [10]. Gradient Boosting has recently demonstrated state-of-the-art performance in tasks like readmission prediction [15]. Classification is found in numerous use cases, like triage, no-show prediction, and medication adherence forecasting.

4.2 Clustering

An unsupervised learning method that groups records together based on feature similarity. K-Means is the most used algorithm used to discover clusters like patient phenotypes, behavioral cohorts, service usage patterns etc. More advanced algorithms like DBSCAN, Latent Class Analysis and Gaussian Mixture Models are better at handling non-linear and non-spherical clusters, which is more appropriate for real-world healthcare data. Clustering is

especially useful in tasks that focus on preventative care and life-cycle interventions[16].

4.3 Association Rule Mining

This method is utilized to uncover frequently co-occurring events within healthcare data. Algorithms like Apriori and FP-Growth are used to discover symptom-drug or drug-drug correlations[17]. Additionally, these techniques apply to electronic health records and administrative data, for identifying patient behavior patterns and healthcare pathways.

4.4 Regression

Regression is employed when the target variable is continuous. Linear regression is often used as a baseline, in areas like healthcare cost modeling or hospital length of stay prediction[18]. Methods that improve upon this, like Ridge, Lasso and Elastic Net Regression can handle high-dimensionality data. For more complex relationships between the features, Regression Trees and Random Forest Regressors offer non-linear capabilities. Most regression models can also be used in real-time monitoring and forecasting of resource utilization in hospitals.

4.5 Timeseries Forecasting

In the context of healthcare, temporally ordered data is often available, such as patient monitoring, hospital admissions, disease outbreaks etc. Classic timeseries models like ARIMA and SARIMA are used to forecast trends like ER visits or disease spread. Facebook has developed Prophet, which gained popularity for healthcare data due to its intuitive parameters[19]. For more advanced scenarios, a type of recurrent neural network called Long Short-Term Memory (LSTM) is being used to do sequential patient modeling, especially in the ICU and in remote health monitoring systems.

4.6 Natural Language Processing(NLP)

As mentioned previously, healthcare data is often unstructured, particularly older archives. From doctor notes and discharge papers to patient feedback, NLP enables the extraction of structured information from these documents. TF-IDF and Word2Vec are widely used for vectorizing text. Latent Dirichlet Allocation is useful in topic modeling scenarios. More recently, transformer-fueled models like BERT have pushed the boundaries of contextual understanding and are used to this day for named entity recognition and more [20].

4.7 Dimensionality Reduction - Feature Selection

We mentioned healthcare datasets are often high-dimensional, and dimensionality reduction can help both improve model performance and reduce noise. Principal Component Analysis (PCA) remains the most fundamental method, while others like t-SNE and UMAP are being used for visualizing patient clusters [21]. In predictive modeling, Recursive Feature Elimination (RFE) and L1-regularization are frequently used to select the most information-rich features in huge datasets like genomics or with multi-modal health data

5 Challenges, Limitations, and Future Directions

5.1 Challenges and Limitations

Medical diagnosis is considered as a significant yet obscure task that needs to be carried out precisely and efficiently. The need is for algorithms with very high accuracy because it is an issue of life or death. Within the issue of knowledge integrity assessment, some of the biggest challenges are:

- How to develop efficient algorithms for comparing content of two knowledge versions (before and after).
- How to develop algorithms for evaluating the influence of particular data modifications on statistical importance of individual patterns that are collected with the help of common classes of data mining algorithm.[22]
- The format of data being stored is different in different healthcare organizations. Till date there is no standard format laid down for data storing.
- The quality of data is another very important challenge. We cannot extract meaningful information from data which have no quality. For example datasets with missing values or noisy ones. The heterogeneous and complex nature of healthcare data, encompassing structured records, unstructured notes, and medical images, adds to the complexity of data preprocessing and integration.[23]
- Data sharing is another major challenge. Neither patients nor healthcare organizations are interested in sharing of their private data. Due to this the epidemic situations may get worse, planning to provide better treatments for a large population may not be possible, and difficulty in the detection of fraud and abuse in healthcare insurance companies.
- Another challenge is that in order to build the data warehouse where all the healthcare organizations within a country share their data is very costly and time consuming process.[4]
- The interpretability of data mining models for the healthcare provider is a vital factor. Complicated "black box" models are often poorly understood and thus do not foster trust on the part of clinicians.[24],[25]
- Additionally, one important limitation of medical data is that we cannot apply big data methods without overcoming a few obstacles such as, issues related to security, data structure, patient monitoring, data storage and transferring, issues related to standardization of data.[26]

In summary, applying data mining in the medical field is an incredibly challenging mission due to the idiosyncrasies of the medical profession. It characterizes widespread process that demands thorough understanding of needs of the healthcare organizations. Knowledge gained with the use of techniques of data mining can be used to make successful decisions that will improve success of healthcare organization and health of the patients. Data mining requires suitable technology and analytical techniques, as well as systems for reporting and tracking which can facilitate measuring of results. The available raw medical data are widely distributed, different and voluminous by nature that must be collected and stored in data warehouses in organized form.[27]

5.2 Future Directions

- Pre-processing of data

Handling missing data, is the most time-consuming and costly part of data mining. The most common method used in the papers was deletion or elimination of missing data and outliers. Although, that means that a lot of information is lost during this process. Outliers can be used to gain information about rare forms of diseases[28]. Instead of neglecting the outliers, future research should analyze them to gain insight.

- Automation of data mining process for non-expert users. The end users of data mining in healthcare are doctors, nurses, and healthcare professionals with limited training in analytics. One solution for this problem is to develop an automated (i.e., without human supervision) system for the end users[29]. A cloud-based automated structure to prevent medical errors could also be developed but the task would be challenging as it involves different application areas and one algorithm will not have similar accuracy for all applications[30].
- Prediction error and “The Black Swan” effect. In healthcare, it is better not to predict than making an erroneous prediction[31]. High prediction accuracy for cancer or any other disease does not ensure an accurate application to decision-making. Moreover, prediction models may be better at predicting commonplace events than rare ones. Researchers should develop more sophisticated models to address the unpredictable, “The Black Swan”[32]. Future research should address the implementation challenges of predictive models, especially how the decision-making process should adapt in case of errors and unpredictable incidents.

6 Conclusion

We have shown that taking advantage of Data Mining in the Healthcare sector is a prime evolution in how medical knowledge is extracted, analyzed, and applied. As the amount of data will continue to grow, from electronic health records to diagnostics, treatments and beyond, the need for analytical tools becomes more pressing. Data Mining has been established as a cornerstone in this shift, allowing professionals to interpret and intuit meaningful patterns from raw data, enhancing both operational efficiency and clinical decision-making.

Through the course of this paper, we have explored the most prominent concepts of Data Mining and its growing importance for healthcare support. We have discussed how the healthcare sector is rich in data but simultaneously faces serious challenges such as inefficiencies, high cost and variability in the quality of provided care. Data Mining offers to ameliorate such issues by strengthening and streamlining decision processes, thus reducing medical errors, detecting fraud, personalizing treatment plans and more. Whether it is predicting disease outcomes or better managing hospital resources, the value of this novel, data-driven approach is evident already.

Our research into applications of Data Mining in Healthcare support demonstrated how it enables early disease detection, hospital length of stay forecasting, uncovering drug-drug interactions and supporting the upcoming field of personalization in medicinal care. By taking advantage of Machine Learning, pattern recognition and predictive modeling, healthcare specialists can be more accurate with their diagnoses and better tailor treatment to their patients, ultimately leading to better outcomes and higher satisfaction.

We also examined the most widely used data mining methodologies used in the domain of healthcare, from classification and clustering to regression, association rule mining, and natural language processing. Each one of these techniques has a unique role to play, depending on the specific task, from patient risk stratification and behavioral analysis to forecasting resource utilization. Model performance can be further enhanced using feature selection and dimensionality reduction methods, when dealing with high-dimensional and multi-modal data.

However, many challenges remain. The usual problems of data heterogeneity, lack of standardization, privacy concerns and variable data quality, continuously block integration and widespread adaptation of data driven strategies. What is more, ethical considerations with regards to patient privacy and confidentiality, together with questions about algorithmic transparency make navigating the landscape especially challenging. Despite everything, ongoing research and development are slowly but surely addressing these limitations, promising more trustworthy and impactful future solutions.

To summarize, data mining stands as a vital enabler in the evolution of modern healthcare. Its capacity to unlock actionable insights from complex data ecosystems is not just enhancing care delivery today, but also laying the groundwork for a more intelligent, responsive and patient-centric future in medicine.

References

- [1] Mohammad Hossein Tekieh and Bijan Raahemi. Importance of data mining in healthcare: a survey. In *Proceedings of the 2015 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining 2015*, pages 1057–1062, 2015.
- [2] Neesha Jothi, Nur’Aini Abdul Rashid, and Wahidah Husain. Data mining in healthcare—a review. *Procedia computer science*, 72:306–313, 2015.
- [3] Md Saiful Islam, Md Mahmudul Hasan, Xiaoyi Wang, Hayley D Germack, and Md Noor-E-Alam. A systematic review on healthcare analytics: application and theoretical perspective of data mining. In *Healthcare*, volume 6, page 54. MDPI, 2018.
- [4] Parvez Ahmad, Saqib Qamar, and Syed Qasim Afser Rizvi. Techniques of data mining in healthcare: a review. *International Journal of Computer Applications*, 120(15), 2015.
- [5] Matthew Herland, Taghi M Khoshgoftaar, and Randall Wald. A review of data mining using big data in health informatics. *Journal of Big data*, 1:1–35, 2014.
- [6] IO Ogundele, OL Popoola, OO Oyesola, and KT Orija. A review on data mining in healthcare. *International Journal of Advanced Research in Computer Engineering and Technology (IJARCET)*, 7:698–704, 2018.
- [7] Jionglin Wu, Jason Roy, and Walter F Stewart. Prediction modeling using ehr data: challenges, strategies, and a comparison of machine learning approaches. *Medical care*, 48(6):S106–S113, 2010.
- [8] Konstantina Kourou, Themis P Exarchos, Konstantinos P Exarchos, Michalis V Karamouzis, and Dimitrios I Fotiadis.

- Machine learning applications in cancer prognosis and prediction. *Computational and structural biotechnology journal*, 13:8–17, 2015.
- [9] Rachda Naila Mekhaldi, Patrice Caulier, Sondes Chaabane, Abdelahad Chraibi, and Sylvain Piechowiak. Using machine learning models to predict the length of stay in a hospital setting. In *World conference on information systems and technologies*, pages 202–211. Springer, 2020.
 - [10] Ernest Yeboah Boateng, Daniel A Abaye, et al. A review of the logistic regression model with emphasis on medical research. *Journal of data analysis and information processing*, 7(04):190, 2019.
 - [11] Behrouz Bokharaeian, Alberto Diaz, and Hamidreza Chitsaz. Enhancing extraction of drug-drug interaction from literature using neutral candidates, negation, and clause dependency. *PLoS One*, 11(10):e0163480, 2016.
 - [12] Ke Han, Peigang Cao, Yu Wang, Fang Xie, Jiaqi Ma, Mengyao Yu, Jianchun Wang, Yaoqun Xu, Yu Zhang, and Jie Wan. A review of approaches for predicting drug–drug interactions based on machine learning. *Frontiers in pharmacology*, 12:814858, 2022.
 - [13] Sreya Vadapalli, Habiba Abdelhalim, Saman Zeeshan, and Zeeshan Ahmed. Artificial intelligence and machine learning approaches using gene expression and variant data for personalized medicine. *Briefings in bioinformatics*, 23(5):bbac191, 2022.
 - [14] Sushen Zhang, Seyed Mojtaba Hosseini Bamakan, Qiang Qu, and Sha Li. Learning for personalized medicine: a comprehensive review from a deep learning perspective. *IEEE reviews in biomedical engineering*, 12:194–208, 2018.
 - [15] Nayara Cristina da Silva, Marcelo Keese Albertini, André Ricardo Backes, and Geórgia das Graças Pena. Machine learning for hospital readmission prediction in pediatric population. *Computer Methods and Programs in Biomedicine*, 244:107980, 2024.
 - [16] Richard W Grant, Jodi McCloskey, Meghan Hatfield, Connie Uratsu, James D Ralston, Elizabeth Bayliss, and Chris J Kennedy. Use of latent class analysis and k-means clustering to identify complex patient profiles. *JAMA network open*, 3(12):e2029068–e2029068, 2020.
 - [17] Kanza Gulzar, Muhammad Ayoob Memon, Syed Muhammad Mohsin, Sheraz Aslam, Syed Muhammad Abrar Akber, and Muhammad Asghar Nadeem. An efficient healthcare data mining approach using apriori algorithm: A case study of eye disorders in young adults. *Information*, 14(4):203, 2023.
 - [18] Ahmed I Taloba, Rasha M Abd El-Aziz, Huda M Alshanbari, and Abdal-Aziz H El-Bagoury. Estimation and prediction of hospitalization and medical care costs using regression in machine learning. *Journal of Healthcare Engineering*, 2022(1):7969220, 2022.
 - [19] Mashael Khayyat, Kaouther Laabidi, Nada Almalki, and Maysoun Al-Zahrani. Time series facebook prophet model and python for covid-19 outbreak prediction. *Computers, Materials & Continua*, 67(3), 2021.
 - [20] Alexander Turchin, Stanislav Masharsky, and Marinka Zitnik. Comparison of bert implementations for natural language processing of narrative medical documents. *Informatix in Medicine Unlocked*, 36:101139, 2023.
 - [21] Carla Lázaro and Cecilio Angulo. Using umap for partially synthetic healthcare tabular data generation and validation. *Sensors*, 24(23):7843, 2024.
 - [22] Prakash Mahindrakar and M Hanumanthappa. Data mining in healthcare: A survey of techniques and algorithms with its limitations and challenges. *Int. J. Eng. Res. Appl*, 3(6):937–941, 2013.
 - [23] Ashwin Belle, Raghuram Thiagarajan, SM Reza Soroushmehr, Fatemeh Navidi, Daniel A Beard, and Kayvan Najarian. Big data analytics in healthcare. *BioMed research international*, 2015(1):370194, 2015.
 - [24] Abhimanyu S Ahuja. The impact of artificial intelligence in medicine on the future role of the physician. *PeerJ*, 7:e7702, 2019.
 - [25] Travis B Murdoch and Allan S Detsky. The inevitable application of big data to health care. *Jama*, 309(13):1351–1352, 2013.
 - [26] Laura Elezabeth, Ved P Mishra, and Joanita Dsouza. The role of big data mining in healthcare applications. In *2018 7th International Conference on Reliability, Infocom Technologies and Optimization (Trends and Future Directions)(ICRITO)*, pages 256–260. IEEE, 2018.
 - [27] Shubpreet Kaur and RK Bawa. Future trends of data mining in predicting the various diseases in medical healthcare system. *International Journal of Energy, Information and Communications*, 6(4):17–34, 2015.
 - [28] Dario Antonelli, Elena Baralis, Giulia Bruno, Tania Cerquitelli, Silvia Chiusano, and Naeem Mahoto. Analysis of diabetic patients through their examination history. *Expert Systems with Applications*, 40(11):4672–4678, 2013.
 - [29] Ricardo S Santos, Suzana MF Malheiros, Sergio Cavaleiro, and JM Parente De Oliveira. A data mining system for providing analytical information on brain tumors to public health decision makers. *Computer methods and programs in biomedicine*, 109(3):269–282, 2013.
 - [30] Chia-Ping Shen, Chinburen Jigjidsuren, Sarangerel Dorjgochoo, Chi-Huang Chen, Wei-Hsin Chen, Chih-Kuo Hsu, Jin-Ming Wu, Chih-Wen Hsueh, Mei-Shu Lai, Ching-Ting Tan, et al. A data-mining framework for transnational healthcare system. *Journal of medical systems*, 36:2565–2575, 2012.
 - [31] Andrew Kusiak, Christopher A Caldarone, Michael D Kelleher, Fred S Lamb, Thomas J Persoon, and Alex Burns. Hypoplastic left heart syndrome: Knowledge discovery with a data mining approach. *Computers in Biology and Medicine*, 36(1):21–40, 2006.
 - [32] Nassim Nicholas Taleb. *The Black Swan:: The Impact of the Highly Improbable: With a new section: "On Robustness and Fragility"*, volume 2. Random house trade paperbacks, 2010.